

STATS200 class notes

Erich Trieschman

2021 Fall quarter

Contents

1	Review: Combinatorics and probability	2
1.1	Calculus cheat sheet	2
1.2	Events and sets	3
1.3	Probability	3
1.3.1	Conditional probability	3
1.3.2	Independence	3
2	Random variables and expectation	4
2.1	Expected value	4
2.2	Variance	4
2.3	Covariance	5
2.4	Correlation	5
2.5	Key theorems	5
2.5.1	Iterated expectation	5
2.5.2	Variance decomposition	5
2.5.3	Cauchy-Schwartz inequality	6
2.5.4	Transformations of random variables	6
2.5.5	Jensen inequality	6
2.5.6	Markov inequality	6
2.5.7	Chebyshev inequality	6
2.6	Moment generating function	6
2.6.1	Common MGF derivations	7
3	Common distribution functions	7
3.1	Discrete distribution functions	7
3.1.1	Bernoulli	7
3.1.2	Binomial distribution	7
3.1.3	Geometric distribution	7
3.1.4	Negative binomial	7
3.1.5	Poisson distribution	7
3.2	Continuous distribution functions	8
3.2.1	Uniform distribution	8
3.2.2	Normal distribution	8
3.2.3	Exponential distribution	8
3.2.4	Gamma distribution	8
3.2.5	Cauchy distribution	9
3.2.6	Beta distribution	9
4	Properties of distributions	9
4.1	Joint distributions	9
4.1.1	Distribution of $X + Y$	10
4.1.2	Expectation of joint distributions	10
4.2	Marginal distributions	10
4.3	Conditional distributions	10

5	Convergence and limit theorems	10
5.1	Convergence in probability	10
5.2	Convergence in L_p	11
5.3	Convergence in distribution	11
5.3.1	Convergence in probability \implies convergence in distribution	11
5.3.2	Slutsky's theorem	11
5.3.3	Student's t distribution (example use case of Slutsky)	11
5.4	Law of large numbers	11
5.5	Central limit theorem	12
5.6	Delta method	12
6	Estimation	12
6.1	Mean Squared Error	13
6.2	Method of Moments estimator	13
6.3	Maximum likelihood estimator	13
6.4	Fisher Information	14
6.4.1	Properties of Fischer Information	14
6.4.2	The "Big" theorem: Asymptotic distribution using Fischer Information	15
6.5	Bayes estimator	15
6.5.1	Example Bayes estimator method	15
6.5.2	Bayes estimator properties	15
6.6	Sufficiency	16
6.6.1	Fischer's Factorization Theorem	16
6.6.2	Rao-Blackwell Theorem	16
7	Hypothesis testing	16
7.1	Likelihood ratio	16
7.2	Neyman-Pearson lemma	17
7.3	Uniformly Most powerful test (UMP)	17
7.3.1	Example LR and UMP test	17
7.4	P-values	17
7.5	Generalized Likelihood Ratio test	18
7.5.1	GLR example I	18
7.5.2	GLR example II	18
7.5.3	Testing multinomial distributions	18
8	Helpful applied methods	19
8.1	Derivation of Linear regression	19
8.2	Confidence intervals	19
8.3	Asymptotic distribution of sample variance	19
8.4	Comparing two samples	19
8.5	Fischer's Exact test	19

1 Review: Combinatorics and probability

1.1 Calculus cheat sheet

Logs: $\log_b(M * N) = \log_b M + \log_b N$ • $\log_b(\frac{M}{N}) = \log_b M - \log_b N$ • $\log_b(M^k) = k \log_b M$ • $e^n e^m = e^{n+m}$

Derivatives: $(x^n)' = nx^{n-1}$ • $(e^x)' = e^x$ • $(e^{u(x)})' = u'(x)e^x$ • $(\log_e(x))' = (\ln x)' = \frac{1}{x}$ • $(f(g(x)))' = f'(g(x))g'(x)$

Integrals: $\int_a^b f(x)dx = \int_{g(a)}^{g(b)} f(g(u))g'(u)du$ where $g(u) = x$ • $\int_a^b u(x)v'(x)dx = u(b)v(b) - u(a)v(a) - \int_a^b u'(x)v(x)dx$

Infinite series and sums: $e^x = 1 + x + \frac{x^2}{2!} + \frac{x^3}{3!} + \dots = \sum_{n=0}^{\infty} \frac{x^n}{n!}$ • $(1 + \frac{a}{n})^n \longrightarrow e^a$

$\ln(1+x) = 1 - x + \frac{x^2}{2} - \frac{x^3}{3} + \dots = \sum_{n=0}^{\infty} (-1)^n \frac{x^n}{n+1}$ • $\frac{1}{1-x} = 1 + x + x^2 + x^3 + \dots = \sum_{n=0}^{\infty} x^n$ for $|x| < 1$

1.2 Events and sets

Set operations follow commutative, associative, and distributive laws:

- Commutative: $E \cup F = F \cup E$ and $E \cap F = F \cap E$ (also written $EF = FE$)
- Associative: $(E \cup F) \cup G = E \cup (F \cup G)$ and $(E \cap F) \cap G = E \cap (F \cap G)$
- Distributive: $(E \cup F) \cap G = (E \cap G) \cup (F \cap G) = E \cap G \cup F \cap G$ and $E \cap F \cup G = (E \cup G) \cap (F \cup G) = E \cup G \cap F \cup G$

DeMorgan's Laws relate the complement of a union to the intersection of complements:

$$(\cup_{i=1}^n E_i)^c = \cap_{i=1}^n E_i^c \bullet (\cap_{i=1}^n E_i)^c = \cup_{i=1}^n E_i^c$$

1.3 Probability

A **probability space** is defined by a triple of objects (S, \mathcal{E}, P) :

- S : Sample space
- \mathcal{E} : Set of possible events within the sample space. Set of events are assumed to be θ -field (below)
- P : Probability for each event

A **θ -field** is a collection of subsets $\mathcal{E} \subset S$ that satisfy $0 \in \mathcal{E} \bullet E \in \mathcal{E} \Rightarrow E^C \in \mathcal{E} \bullet E_i \in \mathcal{E}$ for $1, 2, \dots \Rightarrow \cup_{i=1}^\infty E_i \in \mathcal{E}$

Probability properties:

$$P(A^C) = 1 - P(A) \bullet P(0) = 0 \bullet A \subset B \longrightarrow P(A) \leq P(B) \bullet P(A \cup B) = P(A) + P(B) - P(A \cap B)$$

The **law of total probability** relates marginal probabilities to conditional probabilities. For a partition, E_1, E_2, \dots of set, S , where a partition implies i) E_i, E_j are pairwise disjoint and ii) $\cup_{i=1}^\infty E_i = S$, then

$$P(A) = \sum_{i=1}^\infty P(A \cap E_i) = \sum_{i=1}^\infty P(A | E_i)P(E_i)$$

The **continuity of probability measures** state

- (i) $E_1 \subset E_2 \subset \dots$ Let $E_\infty = \cup_i E_i$, then $P(E_n) \longrightarrow P(E_\infty)$ as $n \longrightarrow \infty$
- (ii) $E_1 \supset E_2 \supset \dots$ Let $E_\infty = \cap_i E_i$, then $P(E_n) \longrightarrow P(E_\infty)$ as $n \longrightarrow \infty$

1.3.1 Conditional probability

The conditional probability is the probability of one event occurring, given the other event occurring. A reframing of conditional probability (see formula below) is the probability of both events occurring, divided by the marginal probability of one of the events occurring.

$$p_{X|Y}(x|y) = \frac{p_{x,y}(x,y)}{p_y(y)}$$

Bayes Theorem leverages conditional probabilities of measured events to glean conditional probabilities of un-measured events:

$$P(E_i | B) = \frac{P(B | E_i)P(E_i)}{\sum_{j=1}^\infty P(B | E_j)P(E_j)} = \frac{P(B | E_i)P(E_i)}{P(B)}$$

Where E_1, E_2, \dots form a partition of the sample space.

1.3.2 Independence

Events A and B are independent if $P(A \cap B) = P(A)P(B)$

It is possible for events to be pairwise independent, but not mutually independent. For example, toss a pair of dice and let D_1 be the number for die 1 and D_2 be the number for die 2. Define $E_i = \{D_i \leq 2\}$. And define $E_3 = \{3 \leq \max(D_1, D_2) \leq 4\}$. These events are pairwise independent, but $P(E_1 \cap E_2 \cap E_3) = 0$, so they are not mutually independent.

2 Random variables and expectation

Random variables are functions connecting a sample space to real numbers: $\{\omega \in S : X(\omega) \leq t\} \in \mathcal{E}$. For example, if coin tosses produce a sample space of {Heads, Tails}, a random variable can be the number of heads.

2.1 Expected value

$$E(X) = \sum_x xP(X = x)$$

Which can also be written as

$$E(X) = \sum_{s \in S} X(s)p(s), \text{ where } p(s) \text{ is the probability that element } s \in S \text{ occurs:}$$

Proof:

$$\begin{aligned} E(X) &= \sum_i x_i P(X = x_i), \text{ for } E_i = \{X = x_i\} = \{s \in S : X(s) = x_i\} \\ &= \sum_i x_i \sum_{s \in E_i} p(s) = \sum_i \sum_{s \in E_i} x_i p(s) = \sum_i \sum_{s \in E_i} X(s) p(s) = \sum_{s \in S} x_i p(s) \end{aligned}$$

This equation structure helps proof several properties of the expected value:

- $E(g(X)) = \sum_i g(x_i) p_X(x_i)$, assuming $g(x_i) = y_i$

$$\sum_i g(x_i) p_X(x_i) = \sum_j \sum_{i: g(x_i) = y_j} g(x_i) p_X(x_i) = \sum_j \sum_{i: g(x_i) = y_j} y_j p_X(x_i) = \sum_j y_j P(g(X) = y_j) = E(g(X))$$

- $E(aX + b) = aE(X) + b$

$$E(aX + b) = \sum_{s \in S} (aX(s) + b)p(s) = a \sum_{s \in S} X(s)p(s) + \sum_{s \in S} bp(s) = aE(X) + b$$

- $E(X + Y) = E(X) + E(Y)$

$$E(X + Y) = \sum_{s \in S} (X(s) + Y(s))p(s) = \sum_{s \in S} X(s)p(s) + \sum_{s \in S} Y(s)p(s) = E(X) + E(Y)$$

2.2 Variance

$$Var(X) = E((X - E(X))^2) = \sigma^2 \bullet SD = \sqrt{Var(X)} = \sqrt{\sigma^2} = \sigma$$

Several properties of variance follow from linearity of expectation:

$$(i) \quad Var(X) = E(X^2) - \mu^2$$

$$Var(X) = E((X - \mu)^2) = E(X^2 - 2X\mu + \mu^2) = E(X^2 - 2\mu X + \mu^2) = E(X^2) - 2\mu^2 + \mu^2 = E(X^2) - \mu^2$$

$$(ii) \quad Var(aX + b) = a^2 Var(X)$$

$$Var(aX + b) = E((aX + b)^2) - E(aX + b)^2 = E(a^2 X^2 + 2abX + b^2) - (aE(X) + b)^2$$

$$Var(aX + b) = a^2 E(X^2) + 2abE(X) + b^2 - a^2 E(X)^2 - 2abE(X) - b^2 = a^2 E(X^2) - a^2 E(X)^2 = a^2 (E(X^2) - E(X)^2)$$

$$(iii) \quad Var(X + Y) = Var(X) + Var(Y) \text{ for } X, Y \text{ independent}$$

$$Var(X + Y) = E((X + Y)^2) - E(X + Y)^2 = E(X^2) + 2E(XY) + E(Y^2) - E(X^2) - 2E(X)E(Y) - E(Y)^2$$

$$Var(X + Y) = E(X^2) - E(X)^2 + E(Y^2) - E(Y)^2, \text{ since } E(XY) = 0 \text{ (by independence) and } E(X) = E(Y) = 0 \text{ (WLOG)}$$

$$Var(X + Y) = Var(X) + Var(Y)$$

2.3 Covariance

$$Cov(X, Y) = E((X - E(X))(Y - E(Y))) = E(XY) - E(X)E(Y)$$

Several properties of covariance follow from linearity of expectation

$$(i) Cov(X, X) = Var(X) :$$

$$Cov(X, X) = E[(X - E(X))(X - E(X))] = E[(X - E(X))^2] = Var(X)$$

$$(ii) Cov(X, Y) = E(XY) - E(X)E(Y) :$$

$$Cov(X, Y) = E[(X - E(X))(Y - E(Y))] = E(XY - E(Y)X - E(X)Y + E(X)E(Y))$$

$$Cov(X, Y) = E(XY) - E(X)E(Y) - E(X)E(Y) + E(X)E(Y) = E(XY) - E(X)E(Y)$$

$$(iii) \text{ if } X, Y \text{ independent, then } Cov(X, Y) = 0$$

$$(iv) Cov(aX, bY) = abCov(X, Y) :$$

$$Cov(aX, bY) = E(abXY) - E(aX)E(bY) = ab(E(XY) - E(X)E(Y)) = abCov(X, Y)$$

$$(v) Cov(X, Y + Z) = Cov(X, Y) + Cov(X, Z) :$$

$$Cov(X, Y + Z) = E(X(Y + Z)) - E(X)E(Y + Z)$$

$$Cov(X, Y + Z) = E(XY) + E(XZ) - E(X)E(Y) - E(X)E(Z) = Cov(X, Y) + Cov(X, Z)$$

$$(vi) Cov(U, V) = \sum_i \sum_j b_i d_j Cov(X_i, Y_j), \text{ with } U = a + \sum_i b_i X_i \text{ and } V = c + \sum_j d_j Y_j :$$

$$(vii) Var(X + Y) = Var(X) + Var(Y) + 2Cov(X, Y) :$$

$$Var(X + Y) = Cov(X + Y, X + Y) = Cov(U, V), \text{ for } U = V = X + Y$$

$$Var(X + Y) = Cov(U, V) = Cov(X, X) + Cov(X, Y) + Cov(Y, Y) + Cov(Y, X), \text{ using vi}$$

$$Var(X + Y) = Var(X) + Var(Y) + 2Cov(X, Y)$$

2.4 Correlation

$$\rho = \frac{Cov(X, Y)}{\sqrt{Var(X)Var(Y)}}$$

2.5 Key theorems

2.5.1 Iterated expectation

Law of iterated expectation: $E(E(Y | X)) = E(Y)$

Proof:

$$E(Y | X) = \sum_y y \frac{f_{X,Y}(X, y)}{f_X(X)}$$

$$E(E(Y | X)) = \sum_x \sum_y \left(y \frac{f_{X,Y}(x, y)}{f_X(x)} \right) f_X(x) = \sum_x \sum_y y f_{X,Y}(x, y) = \sum_y y f_Y(y) = E(Y)$$

2.5.2 Variance decomposition

Variance decomposition formula: $Var(Y) = E(Var(Y | X)) + Var(E(Y | X))$

2.5.3 Cauchy-Schwartz inequality

Cauchy-Schwartz inequality: $E(UV)^2 \leq E(U^2)E(V^2)$, with equality if $P(cU = V) = 1$ for some constant, c
Proof:

$$\text{let } h(t) = E((tU - V)^2) \geq 0, \quad h(t) = t^2 E(U^2) - 2tE(UV) + E(V^2), \text{ a quadratic equation} \\ h(t) \geq 0 \Rightarrow \text{discriminant} \leq 0 \iff 4E(UV)^2 - 4E(U^2)E(V^2) \leq 0 \iff E(UV)^2 \leq E(U^2)E(V^2)$$

2.5.4 Transformations of random variables

For X with density f_X and $Y = g(X)$

$$F_Y(y) = P(g(X) \leq y) = P(X \leq g^{-1}(y)) = F_X(g^{-1}(y)) \\ f_Y(y) = \frac{d}{dy} F_X(g^{-1}(y)) = f_X(g^{-1}(y)) \left| \frac{d}{dy} g^{-1}(y) \right|$$

Note: When computing $F_X(g^{-1}(y))$ be wary of how sign changes may affect the inequality.

2.5.5 Jensen inequality

Jensen inequality: $E(g(x)) \geq g(E(x))$ for $g(x)$ convex

Proof: Let $E(X) = \mu$, and $L(X)$ a line s.t. $L(\mu) = g(E(x))$:

$$g(X) \geq L(X) \text{ for all } X \iff E(g(X)) \geq E(L(X)) = L(E(X)) = g(E(X))$$

2.5.6 Markov inequality

Markov inequality: For $X \geq 0$, $P(X \geq t) \leq \frac{E(X)}{t} \quad \forall t > 0$

Proof:

$$\text{Let } Y = \begin{cases} 1 & X \geq t \\ 0 & \text{otherwise} \end{cases}, \text{ Then } tY \leq X \text{ since } \begin{cases} X \geq t & t * 1 \leq X \\ X < t & t * 0 < X \end{cases} \\ tY \leq X \implies E(tY) \leq E(X) \implies tP(X \geq t) \leq E(X) \implies P(X \geq t) \leq \frac{E(X)}{t}$$

2.5.7 Chebyshev inequality

Chebyshev inequality: $P(|X - E(X)| \geq t) \leq \frac{Var(X)}{t^2} \quad \forall t > 0$

Proof:

$$P(|X - E(X)| \geq t) = P((X - E(X))^2 \geq t^2) \leq \frac{E((X - E(X))^2)}{t^2}, \text{ by Markov inequality} \\ P((X - E(X))^2 \geq t^2) \leq \frac{Var(X)}{t^2}$$

2.6 Moment generating function

The moment generating function of a random variable X is defined as

$$M_X(t) = \mathbb{E}[e^{tX}] = \sum_{n=0}^{\infty} \frac{\mathbb{E}[X^n]}{n!} t^n \leftarrow \text{power series}$$

Notice its called a moment generating function because each derivative of this function can generate a new moment of X at $t = 0$:

$$M_X^{(n)}(0) = \mathbb{E}[X^n]$$

2.6.1 Common MGF derivations

- $Y = a + bX \implies M_Y = e^{at} M_X(bt)$
- $Z = X + Y, X \perp Y \implies M_Z = M_Y M_X = E(e^t X) E(e^t Y)$

3 Common distribution functions

3.1 Discrete distribution functions

3.1.1 Bernoulli

Probability mass function (*Bernouli*(p)): Random variable X takes the value 1 with probability p and the value 0 with probability $1 - p$

$$p(x) = p^x (1 - p)^{1-x}, \quad x \in \{0, 1\}$$

Expected value: p

Variance: $p(1 - p)$

3.1.2 Binomial distribution

Probability mass function (*Bin*(n, p)): For random variable X , the number of successes in n trials, the probability of observing j successes where each success has probability p is

$$P(X = j) = \binom{n}{j} p^j (1 - p)^{n-j}$$

Expected value: np

Variance: $np(1 - p)$

MLE: $\hat{p} = X/n$

3.1.3 Geometric distribution

Probability mass function (*Geom*(p)): For random variable X , the number of trials until the first success (included) with probability p is

$$P(X = j) = (1 - p)^{j-1} p$$

Expected value: $\frac{1}{p}$

Variance: $\frac{1-p}{p^2}$

3.1.4 Negative binomial

Probability mass function (*NB*(r, p)): For random variable X , the number of successes, k before a specified number of failures, r , with probability of success p is

$$P(X = k) = \binom{k+r-1}{k} (1-p)^r p^k$$

Expected value: $\frac{pr}{1-p}$

Variance: $\frac{pr}{(1-p)^2}$

3.1.5 Poisson distribution

Probability mass function (*Pois*(λ)): For random variable, X , the number of events, k , occurring in a fixed interval of time or space if these events occur with a known constant mean rate, λ

$$P(X = k) = \frac{\lambda^k e^{-\lambda}}{k!}$$

Expected value: λ

Variance: λ

MLE: $\hat{\lambda} = \bar{X}$

- $X_1, \dots, X_n \stackrel{i.i.d}{\sim} \text{Poisson}(\lambda_i) \implies \sum_{i=1}^n X_i \sim \text{Poisson}(\sum_{i=1}^n \lambda_i)$

3.2 Continuous distribution functions

3.2.1 Uniform distribution

Unif(a, b): The distribution describes an experiment where there is an arbitrary outcome that lies between certain bounds.[1] The bounds are defined by the parameters, a and b , which are the minimum and maximum values

$$pdf : f(x) = \begin{cases} \frac{1}{b-a} & \text{for } x \in [a, b] \\ 0 & \text{otherwise} \end{cases} \bullet cdf : F(x) = \begin{cases} 0 & \text{for } x < a \\ \frac{x-a}{b-a} & \text{for } x \in [a, b] \\ 1 & \text{for } x > b \end{cases}$$

Expected value: $\frac{1}{2}(a+b)$

Variance: $\frac{1}{12}(b-a)^2$

MLE: $\hat{\theta} = X_{(n)} = \max\{X_1, \dots, X_n\}$

3.2.2 Normal distribution

$N(\mu, \sigma)$

$$pdf : f(x) = \frac{1}{\sigma\sqrt{2\pi}} \exp\left(-\frac{(x-\mu)^2}{2\sigma^2}\right) \bullet cdf : F(x) = \frac{1}{\sqrt{2\pi}} \int_{-\infty}^x e^{-t^2/2} dt$$

Expected value: μ

Variance: σ^2

MLE: $\hat{\mu} = \bar{X} \bullet \hat{\sigma}^2 = \frac{1}{n} \sum_{i=1}^n (X_i - \bar{X})^2$

$$\bullet X_i \sim N(0, 1) \implies \sum_{i=1}^n X_i \sim N(0, n) \implies \frac{1}{n} \sum_{i=1}^n X_i \sim N(0, n/n^2) = N(0, 1/n)$$

$$\bullet \frac{(\bar{Y}_m - \bar{X}_n) - (\mu_Y - \mu_X)}{\sqrt{\sigma_X^2/n + \sigma_Y^2/m}} \sim Z = N(0, 1)$$

3.2.3 Exponential distribution

Exp(λ): the probability distribution of the time between events in a Poisson point process, i.e., a process in which events occur continuously and independently at a constant average rate. It is a particular case of the gamma distribution.

$$pdf : f(x) = \lambda e^{-\lambda x} \bullet cdf : F(x) = 1 - e^{-\lambda x}$$

Expected value: $\frac{1}{\lambda}$

Variance: $\frac{1}{\lambda^2}$

MLE: $\hat{\lambda} = 1/\bar{X}$

3.2.4 Gamma distribution

Gamma(α, λ): a two-parameter family of continuous probability distributions.

$$pdf : f(x) = \frac{\lambda^\alpha}{\Gamma(\alpha)} x^{\alpha-1} e^{-\lambda x}, \text{ where } \Gamma(\alpha) = (\alpha-1)! \text{ for any positive integer, } \alpha$$

$$cdf : F(x) = \frac{1}{\Gamma(\alpha)} \gamma(\alpha, \lambda x), \text{ where } \gamma(\alpha, x) = \int_0^x t^{\alpha-1} e^{-t} dt$$

Expected value: $\frac{\alpha}{\lambda}$

Variance: $\frac{\alpha}{\lambda^2}$

3.2.5 Cauchy distribution

Cauchy(t, s): The Cauchy distribution is often used in statistics as the canonical example of a "pathological" distribution since both its expected value and its variance are undefined

$$\begin{aligned} pdf : f(x) &= \frac{1}{s\pi(1 + (x - t)/s)^2}, \text{ where } s \text{ is the scale parameter and } t \text{ is the location parameter} \\ cdf : &\frac{1}{\pi} \arctan\left(\frac{x - t}{s}\right) + \frac{1}{2} \end{aligned}$$

Expected value: *DNE*

Variance: *DNE*

3.2.6 Beta distribution

Beta(α, β): a family of continuous probability distributions defined on the interval $[0, 1]$ parameterized by two positive shape parameters that appear as exponents of the random variable and control the shape of the distribution.

$$pdf : f(x) = \frac{\Gamma(\alpha + \beta)}{\Gamma(\alpha)\Gamma(\beta)} x^{\alpha-1} (1 - x)^{\beta-1}, \text{ where } x \in [0, 1], \text{ and } \Gamma(k) = (k - 1)! \text{ for any positive integer } k$$

Expected value: $\frac{\alpha}{\alpha + \beta}$

Variance: $\frac{\alpha\beta}{(\alpha + \beta)^2(\alpha + \beta + 1)}$

4 Properties of distributions

Distributions provide a relationship between a random variable and the probability of observing that random variable.

4.1 Joint distributions

General case:

$$\begin{aligned} \text{cdf: } F_{X_1, \dots, X_n}(x_1, \dots, x_n) &= P(X_1 \leq x_1, \dots, X_n \leq x_n) \iff P((X_1, \dots, X_n) \in E) = \int \dots \int_E f_{X_1, \dots, X_n} dx_1 \dots dx_n \\ \text{pmf: } f_{X_1, \dots, X_n}(x_1, \dots, x_n) &= P(X_1 = x_1, \dots, X_n = x_n) \end{aligned}$$

When X_i independent:

$$\begin{aligned} \text{cdf: } P(X_1 \leq x_1, \dots, X_n \leq x_n) &= P(X_1 \leq x_1) \dots P(X_n \leq x_n) = \prod_{i=1}^n P(X_i \leq x_i) \\ \text{pmf: } P(X_1 = x_1, \dots, X_n = x_n) &= P(X_1 = x_1) \dots P(X_n = x_n) = \prod_{i=1}^n P(X_i = x_i) \end{aligned}$$

4.1.1 Distribution of $X + Y$

The distribution of a sum of random variables is called a **convolution**. For X, Y independent

$$\begin{aligned}
 F_{X+Y}(t) &= P(X + Y \leq t) = P(X \leq t - y) \\
 &= \int_{-\infty}^{\infty} P(X \leq t - y \mid Y = y) f_Y(y) dy, \text{ to get marginal distribution} \\
 &= \int_{-\infty}^{\infty} F_X(t - y) f_Y(y) dy, \text{ since } X, Y \text{ independent} \\
 f_{X+Y}(t) &= \int_{-\infty}^{\infty} f_X(t - y) f_Y(y) dy \\
 p_{X+Y}(t) &= P(X + Y = t) = \sum_{x=-\infty}^{\infty} p_X(t - y) p_Y(y)
 \end{aligned}$$

4.1.2 Expectation of joint distributions

For X, Y joint distribution, $f_{X,Y}(x, y)$, or probability mass function, $p(x, y)$

$$\begin{aligned}
 \text{pmf: } E[g(X, Y)] &= \sum_s g(X(s), Y(s)) p(s) = \sum_x \sum_y g(x, y) \sum_{s: X(s)=x, Y(s)=y} p(s) = \sum_x \sum_y g(x, y) p(x, y) \\
 \text{pdf: } E[g(X, Y)] &= \int_{y=-\infty}^{\infty} \int_{x=-\infty}^{\infty} g(x, y) f(x, y) dx dy
 \end{aligned}$$

4.2 Marginal distributions

Marginal density functions or marginal probability mass functions are obtained by integrating or summing out the other variables

$$\text{pmf: } p_Y(y) = \sum_x y P(Y = y \mid x) \bullet \text{pdf: } F_Y(y) = \int_a^b f(x, y) dx, \text{ where } x \in [a, b]$$

4.3 Conditional distributions

Law of total probability:

$$\begin{aligned}
 P(E) &= \sum_{i=-\infty}^{\infty} P(E \mid X = x) P(X) \text{ and } P(E) = \int_{-\infty}^{\infty} P(E \mid X = x) f(x) dx \\
 \text{Recall: } p_{X|Y}(x|y) &= \frac{p_{x,y}(x, y)}{p_Y(y)} \text{ and } f_{X|Y}(x|y) = \frac{f_{X,Y}(x, y)}{f_Y(y)}
 \end{aligned}$$

5 Convergence and limit theorems

5.1 Convergence in probability

A sequence of random variables, X_n , converges in probability, $X_n \xrightarrow{p} X$ when $P(|X_n - X| > \epsilon) \rightarrow 0$ as $n \rightarrow \infty$

Consistent estimator: $T_n = T_n(X_1, \dots, X_n)$ converges in probability to $g(\theta)$, a function of the model parameter

Additional properties of convergence in probability

- if $X_n \xrightarrow{p} X$ and $a_n \xrightarrow{p} a$ then $a_n X_n \xrightarrow{p} aX$
- if $X_n \xrightarrow{p} X$ and $A_n \xrightarrow{p} A$ then $A_n X_n \xrightarrow{p} AX$
- if $X_n \xrightarrow{p} X$, $A_n \xrightarrow{p} A$, and $B_n \xrightarrow{p} B$ then $A_n X_n + B_n \xrightarrow{p} AX + B$
- if $X_n \xrightarrow{p} X$ and g a continuous function then $g(X_n) \xrightarrow{p} g(X)$ (**continuous mapping theorem**)

5.2 Convergence in L_p

See https://en.wikipedia.org/wiki/Lp_space for more information (not much covered in class).

Convergence in L_p is stronger than convergence in probability. **Counter example** to convergence in probability \implies convergence in L_p :

$$\begin{aligned} \text{Let } X_n &= \begin{cases} n & \frac{1}{n} \\ 0 & 1 - \frac{1}{n} \end{cases} \\ X_n &\xrightarrow{p} 0 : P(|X_n - 0| \geq \epsilon) = P(X_n = n) = 1/n \rightarrow 0 \text{ as } n \rightarrow \infty \\ \text{but } E(X_n) &= n \frac{1}{n} + 0(1 - \frac{1}{n}) = 1 \implies \text{no convergence in } L_p \end{aligned}$$

5.3 Convergence in distribution

A sequence of random vectors, X_n , converges in distribution to a random vector, $X_n \xrightarrow{d} X$ when

$$\lim_{n \rightarrow \infty} F_{X_n}(x) = F_X(x) \text{ at all continuity points in } F_X$$

- Convergence in distribution **does not** imply convergence in probability unless convergence in distribution is to a single point
- if $X_n \xrightarrow{d} X$ and g a continuous function then $g(X_n) \xrightarrow{d} g(X)$ (**continuous mapping theorem**)

5.3.1 Convergence in probability \implies convergence in distribution

Let X have cdf, F , with t a continuity point of F

$$\begin{aligned} P(X_n \leq a) &\leq P(X \leq a + \epsilon) + P(|X_n - X| > \epsilon) \text{ by lemma} \\ P(X \leq a - \epsilon) - P(|X_n - X| > \epsilon) &\leq P(X_n \leq a) \leq P(X \leq a + \epsilon) + P(|X_n - X| > \epsilon) \\ F_X(a - \epsilon) &\leq \lim_{n \rightarrow \infty} P(X_n \leq a) \leq F_X(a + \epsilon), \text{ where } F_X(a) = P(X \leq a) \\ &\implies \lim_{n \rightarrow \infty} P(X_n \leq a) = P(X \leq a) \implies \{X_n\} \xrightarrow{d} X \end{aligned}$$

5.3.2 Slutsky's theorem

$A_n X_n + B_n \xrightarrow{d} aX + b$ if $\{X_n\}$ sequence with $X_n \xrightarrow{d} X$, $\{A_n\}$ sequence with $A_n \xrightarrow{d} A$, $\{B_n\}$ sequence with $B_n \xrightarrow{d} b$

5.3.3 Student's t distribution (example use case of Slutsky)

$$\begin{aligned} \frac{\sqrt{n}(\bar{X}_n - \mu)}{\hat{\sigma}} &= \frac{\sqrt{n}(\bar{X}_n - \mu)}{\sigma} \frac{\sigma}{\hat{\sigma}}, \text{ and we know } \frac{\sqrt{n}(\bar{X}_n - \mu)}{\sigma} \xrightarrow{d} N(0, 1) \text{ and } \frac{\sigma}{\hat{\sigma}} \xrightarrow{p} 1 \text{ since } \hat{\sigma} \xrightarrow{p} \sigma \\ \text{So, by Slutsky's theorem, } &\frac{\sqrt{n}(\bar{X}_n - \mu)}{\hat{\sigma}} \xrightarrow{d} N(0, 1) * 1 \end{aligned}$$

This RHS term is referred to as the t-statistic, which follows a Student's t distribution with $n - 1$ degrees of freedom. In practice, if the sample is reasonably sized, it won't make a difference using the Normal distribution instead of the Student's t distribution.

5.4 Law of large numbers

For X_1, X_2, \dots, X_n i.i.d. with $E(X_i) = \mu$, $Var(X_i) = \sigma^2$, $\bar{X}_n = \frac{1}{n} \sum_{i=1}^n X_i$, then for any $\epsilon > 0$

$$P(|\bar{X}_n - \mu| > \epsilon) \rightarrow 0 \text{ as } n \rightarrow \infty$$

Proof:

$$\begin{aligned}\mathbb{E}(\bar{X}_n) &= \frac{1}{n} \sum_{i=1}^n \mathbb{E}(X_i) = \mu \\ \text{Var}(\bar{X}_n) &= \frac{1}{n^2} \sum_{i=1}^n \text{Var}(X_i) = \frac{\sigma^2}{n}, \text{ since } X_i \text{ independent} \\ P(|\bar{X}_n - \mu| > \epsilon) &\leq \frac{\text{Var}(\bar{X}_n)}{\epsilon^2} = \frac{\sigma^2}{n\epsilon^2} \rightarrow 0 \text{ as } n \rightarrow \infty, \text{ by Chebyshev inequality}\end{aligned}$$

5.5 Central limit theorem

Most useful form of CLT, which can be used for approximate methods:

$$\sqrt{n} \frac{(\bar{X}_n - \mu)}{\sigma} \rightarrow N(0, 1) \iff \sqrt{n}(\bar{X}_n - \mu) \rightarrow N(0, \sigma^2)$$

Formal definition: For X_1, X_2, \dots, X_n i.i.d. with $E(X_i) = 0$ (WLOG), $\text{Var}(X_i) = \sigma^2$, c.d.f. F , and MGF, M , (defined in a neighborhood of zero). Then

$$\lim_{n \rightarrow \infty} P\left(\frac{S_n}{\sigma\sqrt{n}} \leq x\right) = \Phi(x), \text{ for } S_n = \sum_{i=1}^n X_i$$

Proof: Let $Z_n = \frac{S_n}{\sigma\sqrt{n}}$. We show the MGF of Z_n tends to the MGF of the standard normal distribution. Since S_n is a sum of independent random variables,

$$M_{S_n}(t) = [M(t)]^n \text{ and } M_{Z_n}(t) = \left[M\left(\frac{t}{\sigma\sqrt{n}}\right)\right]^n$$

Reminder: Taylor series expansion of $M(s) = M(0) + sM'(0) + \frac{1}{2}sM''(0) + \epsilon_s$

$$M\left(\frac{t}{\sigma\sqrt{n}}\right) = 1 + \frac{1}{2}\sigma^2\left(\frac{t}{\sigma\sqrt{n}}\right)^2 + \epsilon_n \text{ with } E(X) = M'(0) = 0, \text{Var}(X) = M''(0) = \sigma^2$$

$$M_{Z_n}(t) = \left(1 + \frac{t^2}{2n} + \epsilon_n\right)^n \rightarrow e^{\frac{t^2}{2}} \text{ as } n \rightarrow \infty, \text{ by the infinite series convergence to } e^a$$

Since $e^{\frac{t^2}{2}}$ is the MGF of the standard normal distribution, we have proven the central limit theorem.

5.6 Delta method

If g is a differentiable function at μ , $\sqrt{n}(g(\bar{X}_n) - g(\mu)) \xrightarrow{d} N(0, g'(\mu)^2\sigma^2)$

Proof: For general g and assuming $E(\bar{X}_n) = \mu$

$$\begin{aligned}g(\bar{X}_n) &\approx g(\mu) + g'(\mu)(\bar{X}_n - \mu) + \frac{1}{2}g''(\mu)(\bar{X}_n - \mu)^2 + \epsilon \text{ (Taylor approximation of } g(\mu)) \\ g(\bar{X}_n) - g(\mu) &\approx g'(\mu)(\bar{X}_n - \mu) + \epsilon \iff \sqrt{n}(g(\bar{X}_n) - g(\mu)) \approx g'(\mu)\sqrt{n}(\bar{X}_n - \mu) + \epsilon \text{ and we know} \\ \sqrt{n}(\bar{X}_n - \mu) &\xrightarrow{d} N(0, \sigma^2) \iff g'(\mu)\sqrt{n}(\bar{X}_n - \mu) \xrightarrow{d} N(0, g'(\mu)^2\sigma^2)\end{aligned}$$

So $\sqrt{n}(g(\bar{X}_n) - g(\mu)) \xrightarrow{d} N(0, g'(\mu)^2\sigma^2)$

Note: if we find that $g'(\mu) = 0$, then repeat this process with the second derivative, $g''(\mu)$.

6 Estimation

Here we use functions of the data ("estimators"), $T(X_1, \dots, X_n)$ to estimate population parameters, θ

6.1 Mean Squared Error

The **Mean Squared Error (MSE)** can be used to evaluate our estimators.

$$\begin{aligned} MSE(T, \theta) &= E_\theta[(T - g(\theta))^2] = E_\theta(T^2) - 2g(\theta)E_\theta(T) + g(\theta)^2 \\ &= Var_\theta(T) + E_\theta(T)^2 + 2g(\theta)E_\theta(T) + g(\theta)^2 = Var_\theta(T) + (E_\theta(T) - g(\theta))^2 \\ &= Var_\theta(T) + Bias_\theta^2(T), \text{ where } Bias_\theta(T) = E_\theta(T) - g(\theta) \end{aligned}$$

Corollary: an unbiased estimator, T , is one such that $E_\theta(T) = g(\theta)$

6.2 Method of Moments estimator

To generate a method of moments estimator

- Calculate a moment with MGF of the assumed distribution. Any moment, k , can be used, but lower moments will typically lead to an estimator distribution with lower variance

$$E(X^k) = g(\theta)$$

- Invert this expression to create an expression for the parameter(s) in terms of the moment

$$g^{-1}(E(X^k)) = \theta \implies f(E(X^k)) = \theta, \text{ where } f(x) = g^{-1}(x)$$

- Insert the sample moment into this expression, thus obtaining estimates of the parameters in terms of data

$$\hat{\theta} = f\left(\frac{1}{n} \sum X_i^k\right), \text{ by LNN } \frac{1}{n} \sum X_i^k \xrightarrow{p} E(X^k)$$

- Use the delta method to determine what the method of moments estimator converges to in distribution

$$\sqrt{n}\left(f\left(\frac{1}{n} \sum X_i^k\right) - \theta\right) \xrightarrow{d} N(0, f'(E(X_i^k))^2 Var(X_i^k)^2)$$

Methods of moment estimators are not uniquely determined, nor must they exist. The motivation for subsequent estimators is to help us pick the estimator with the smallest possible variance.

6.3 Maximum likelihood estimator

The **maximum likelihood estimator** constructs an estimator, $\hat{\theta}_{MLE}$, that maximizes the likelihood function with respect to θ .

The **likelihood function**, $L(\theta)$ is the joint density or probability mass function, $f(X, \theta)$, evaluated at the data, $\{X_i, \dots, X_n\}$. Assuming the data is *i.i.d.*:

$$L(\theta) = \prod_{i=1}^n f(X_i, \theta)$$

General approach to constructing MLE:

- Construct the likelihood function: $L(\theta) = \prod_{i=1}^n f(X_i, \theta)$

$$\text{Example normal: } L(\theta) = \prod_{i=1}^n \frac{1}{\sigma\sqrt{2\pi}} \exp\left(-\frac{(X_i - \mu)^2}{2\sigma^2}\right) = \frac{1}{(\sigma\sqrt{2\pi})^n} \exp\left(-\frac{1}{2\sigma^2} \sum_{i=1}^n (X_i - \mu)^2\right)$$

$$\text{Example restricted multinomial: } L(\theta) \propto f_1(\theta)^{X_1} \dots f_k(\theta)^{X_k}$$

- Take the log of the likelihood: $\log(L(\theta)) = l(\theta) = \sum_{i=1}^n \log(f(X_i, \theta))$
- Take the derivative of the log-likelihood function with respect to θ : $\frac{d}{d\theta} l(\theta) = \sum_{i=1}^n \frac{d}{d\theta} \log(f(X_i, \theta))$

- Find critical points of this function ($0 = \sum_{i=1}^n \frac{d}{d\theta} \log(f(X_i, \hat{\theta}))$) and determine that one is a maximum ($0 = \sum_{i=1}^n \frac{d^2}{d\theta^2} \log(f(X_i, \hat{\theta}))$), checking if $\hat{\theta} < 0$)

Approach to constructing MLE when indicators, $\mathbb{I}\{U\}$, are present:

- Logs of indicators and derivatives of indicators are very difficult to work with
- Simplify likelihood function (splitting indicators when possible)
- Make an argument for why the function is increasing or decreasing
- Determine the value at the bounds of the function

6.4 Fisher Information

The **information** that data, X , contains about parameter, θ is defined by

$$I(\theta) = E_{\theta} \left[\left(\frac{d}{d\theta} \log(f(X, \theta)) \right)^2 \right]$$

Fisher Information assumes **differentiability** and **existence of the second moment**. $\frac{d}{d\theta} \log(f(X, \theta))$ is called the **score** function

6.4.1 Properties of Fischer Information

1. $E_{\theta} \left[\left(\frac{d}{d\theta} \log(f(X, \theta)) \right) \right] = 0$:

$$E_{\theta} \left[\left(\frac{d}{d\theta} \log(f(X, \theta)) \right) \right] = \int \frac{d}{d\theta} \log(f(x, \theta)) f(x, \theta) dx = \int \frac{f'(x, \theta)}{f(x, \theta)} f(x, \theta) dx = \int f'(x, \theta) dx = \frac{d}{d\theta} \int f(x, \theta) dx = \frac{d}{d\theta} * 1 = 0$$

2. $I(\theta) = Var \left(\frac{d}{d\theta} \log(f(X, \theta)) \right)$:

$$Var \left(\frac{d}{d\theta} \log(f(X, \theta)) \right) = E_{\theta} \left[\left(\frac{d}{d\theta} \log(f(X, \theta)) \right)^2 \right] - E_{\theta} \left[\left(\frac{d}{d\theta} \log(f(X, \theta)) \right) \right]^2 = I(\theta) - 0^2 = I(\theta)$$

3. $I(\theta) = -E_{\theta} \left[\frac{d^2}{d\theta^2} \log(f(X, \theta)) \right]$:

$$\frac{d}{d\theta} \log(f(x, \theta)) = \frac{f'(x, \theta)}{f(x, \theta)} \implies \frac{d^2}{d\theta^2} \log(f(x, \theta)) = \frac{f(x, \theta)f''(x, \theta) - f'(x, \theta)^2}{f(x, \theta)^2}$$

$$E \left[\frac{d^2}{d\theta^2} \log(f(x, \theta)) \right] = \int \frac{f(x, \theta)f''(x, \theta) - f'(x, \theta)^2}{f(x, \theta)^2} f(x, \theta) dx = \int f''(x, \theta) - I(\theta) = -I(\theta), \text{ since } \int \frac{d^2}{d\theta^2} f(x, \theta) = \frac{d^2}{d\theta^2} * 1 = 0$$

4. $I_{X,Y}(\theta) = I_X(\theta) + I_Y(\theta)$ for X, Y independent : (Information increases with larger sample!)

Corollary: $I_n(\theta) = nI_1(\theta)$ for X_1, \dots, X_n i.i.d with $I_1(\theta)$ the Information based on one data

5. **Cramer-Rau-Fisher Inequality:** $Var(T(X)) \geq \frac{g'(\theta)^2}{I(\theta)}$ for $E(T(X)) = g(\theta)$:

$$Cov[T(X), \frac{d}{d\theta} \log(f(X, \theta))] = E[T(X) \frac{d}{d\theta} \log(f(X, \theta))], \text{ using property 1}$$

$$Cov[T(X), \frac{d}{d\theta} \log(f(X, \theta))] = \int T(x) f'(x, \theta) dx = \frac{d}{d\theta} \int T(x) f(x, \theta) dx = \frac{d}{d\theta} E(T(X)) = \frac{d}{d\theta} g(\theta) = g'(\theta)$$

$$g'(\theta)^2 \leq Var(T(X)) Var \left(\frac{d}{d\theta} \log(f(X, \theta)) \right) = Var(T(X)) I(\theta) \text{ by correlation inequality: } \rho^2 \leq 1$$

$$Var(T(X)) \geq \frac{g'(\theta)^2}{I(\theta)}$$

6.4.2 The "Big" theorem: Asymptotic distribution using Fischer Information

Under regularity assumptions, the maximum likelihood estimator (or any other reasonable estimator), $\hat{\theta}$ of θ satisfies

$$\sqrt{n}(\hat{\theta} - \theta) \xrightarrow{d} N\left(0, \frac{1}{I(\theta)}\right)$$

Sketch of proof:

$$L(\theta) = \prod_{i=1}^n f(X_i, \theta) \iff l(\theta) = \log(L(\theta)) = \sum_{i=1}^n \log(f(X_i, \theta))$$

MLE solves $l'(\hat{\theta}) = 0$, with $l'(\theta) \approx l'(\theta_0) + (\hat{\theta} - \theta_0)l''(\theta_0)$ (full proof requires showing the error in this approx. is small)

$$0 = l'(\theta_0) + (\hat{\theta} - \theta_0)l''(\theta_0) \implies \hat{\theta} - \theta_0 = \frac{l'(\theta_0)}{l''(\theta_0)} \iff \sqrt{n}(\hat{\theta} - \theta_0) = \sqrt{n} \frac{l'(\theta_0)}{l''(\theta_0)} = \frac{l'(\theta_0)}{\sqrt{n}} \div \frac{l''(\theta_0)}{n}$$

$$\frac{l''(\theta_0)}{n} = \frac{\sum \frac{d^2}{d\theta^2} \log(f(X, \theta))}{n} \xrightarrow{p} -E_{\theta} \left[\frac{d^2}{d\theta^2} \log(f(X, \theta)) \right] = I(\theta)$$

$$\frac{l'(\theta_0)}{\sqrt{n}} = \frac{\sum \frac{d}{d\theta} \log(f(X, \theta))}{\sqrt{n}} \xrightarrow{d} N(0, I(\theta))$$

$$\frac{l'(\theta_0)}{\sqrt{n}} \div \frac{l''(\theta_0)}{n} \xrightarrow{d} N\left(0, \frac{I(\theta)}{I(\theta)^2}\right) = N\left(0, \frac{1}{I(\theta)}\right), \text{ by Slutsky's theorem}$$

Corollary: $Var(\hat{\theta}_{MLE}) = 1/I(\theta)$

6.5 Bayes estimator

- **Prior distribution:** $\pi(\theta)$ the distribution of random variable Θ from which model parameter θ is drawn.
- **Conditional distribution:** $f(\{X_1, \dots, X_n\} | \theta)$ is the conditional distribution of the data given $\Theta = \theta$
- **Posterior distribution:** $\pi(\theta | \{X_1, \dots, X_n\})$ is the density of the random variable Θ given the observed data

$$\pi(\theta | \{X_1, \dots, X_n\}) = \frac{f(\{X_1, \dots, X_n\} | \theta) \pi(\theta)}{m(\{X_1, \dots, X_n\})}, \text{ for } m(\{X_1, \dots, X_n\}) = \int_{-\infty}^{\infty} f(\{X_1, \dots, X_n\} | \theta) \pi(\theta) dx$$

The **Bayes Estimator** is calculated as $E[\pi(\theta | \{X_1, \dots, X_n\})]$.

For recognizable functions, we can back into the posterior distribution. If the prior and conditional distribution are not recognizable, then we use numerical methods, like MCMC to approximate the posterior distribution

6.5.1 Example Bayes estimator method

$$X \sim \text{Poisson}(\theta), \theta \in [0, 1] \quad \pi(\theta) = \exp(\theta)/(e - 1)$$

$$\pi(\theta | X) \propto \frac{\exp(-\theta)\theta^X}{X!} * \frac{\exp(\theta)}{e - 1} \mathbb{I}[\theta \in [0, 1]] \propto \theta^X \mathbb{I}[\theta \in [0, 1]] (\leftarrow \text{with more data, these functions are joint distributions})$$

$$\pi(\theta | X) = (X + 1)\theta^X, \text{ observing } \text{Beta}(x + 1, 1) = \frac{\Gamma(x + 2)}{\Gamma(X + 1)\Gamma(1)}\theta^x = (x + 1)\theta^x, \theta \in [0, 1]$$

$$E[\pi(\theta | X)] = \int_0^1 \theta(X + 1)\theta^X d\theta = \frac{X + 1}{X + 2}$$

6.5.2 Bayes estimator properties

All admissible estimators are Bayes Estimators. An estimator, $T'(\theta)$, is inadmissible if $\exists T$ such that

$$E_{\theta}[(T - \theta)^2] \leq E_{\theta}[(T' - \theta)^2] \forall \theta \text{ and } E_{\theta}[(T - \theta)^2] < E_{\theta}[(T' - \theta)^2] \text{ for some } \theta$$

Absence any data, the Bayes Estimator is the expectation of the prior, $E(\pi(\theta))$

6.6 Sufficiency

A test statistic, $T = T(X_1, \dots, X_n)$ is **sufficient** for θ if $f(X_1, \dots, X_n \mid T = t)$ does not depend on θ

The claim with a sufficient statistic is that there is no loss in throwing away the data as long as you keep the sufficient statistic

6.6.1 Fischer's Factorization Theorem

The **Fischer's Factorization Theorem** states that

$$T(X_1, \dots, X_n) \text{ is sufficient for } \theta \iff \text{joint density } f(X_1, \dots, X_n, \theta) = g(T(X_1, \dots, X_n), \theta)h(X_1, \dots, X_n)$$

Proof: TODO

6.6.2 Rao-Blackwell Theorem

The **Rao-Blackwell Theorem** states

For $\hat{\theta}$ an estimator of θ with $E(\theta) < \infty$ and T sufficient with $\theta^* = E(\theta \mid T)$ then

$$E[(\theta^* - \theta)^2] \leq E[(\hat{\theta} - \theta)^2]$$

Proof: TODO

7 Hypothesis testing

- We assume data, $\{X_1, \dots, X_n\}$ is generated by a distribution with parameter $\theta \in \Omega$ (could be a vector)
- The null hypothesis, H_0 and alternative hypothesis, H_1 , are hypotheses for the true value of θ
 - A simple hypothesis is for a single value of θ , $H_i : \theta = \theta_i$
 - A composite hypothesis is for a range of θ , $H_i : \theta > 1$ or $H_i : \theta \neq \theta_0$
- The goal in testing is to construct a rule to decide whether to reject H_0
 - Want: $P_{H_0}(\text{falsely rejecting } H_0) = P_{H_0}(\text{Type I error}) \leq \alpha$
 - Want: maximal $P_{H_1}(\text{correctly rejecting } H_0) = 1 - P_{H_1}(\text{falsely accepting } H_0) = 1 - P_{H_1}(\text{Type II error})$
 - The rejection region, R , can be chosen to maximize correct rejections, subject to a Type I error constraint

7.1 Likelihood ratio

For simple hypotheses, the **Likelihood Ratio** is the ratio of the likelihoods under the alternative and null hypotheses. This ratio helps us boost correct rejections while limiting false rejections.

$$LR = \frac{f_{h_1}(\{X_1, \dots, X_n\})}{f_{h_0}(\{X_1, \dots, X_n\})}$$

We can define our rejection region, R using this the likelihood ratio. Specifically

$$R = \left\{ X : \frac{f_{h_1}(X)}{f_{h_0}(X)} \geq c \right\}$$

And constrain Type I error to level α by solving for c

$$P_{H_0}(\text{Type I error}) = P_{H_0}(R) = P_{H_0} \left(\frac{f_{h_1}(X)}{f_{h_0}(X)} \geq c \right) = \alpha$$

Our power then becomes $P_{H_1}(R)$

7.2 Neyman-Pearson lemma

For *simple hypotheses*, H_0, H_1 , the **Neyman-Pearson lemma** states that the **Likelihood Ratio** level- α test, which rejects H_0 when $LR \geq c$, maximizes power, $P_{H_1}(LR \geq c)$. Any other level- α test, R' , has $P_{H_1}(R') \leq P_{H_1}(LR \geq c)$
Proof:

Let $\phi(x) = \{1 \text{ if } x \in R; 0 \text{ otherwise}\}$, $\phi'(x) = \{1 \text{ if } x \in R'; 0 \text{ otherwise}\}$

Let $S^+ = \{x : \phi(x) = 1, \phi'(x) = 0\}$, $S^- = \{x : \phi(x) = 0, \phi'(x) = 1\}$

$$\begin{aligned} \int_{-\infty}^{\infty} (\phi(x) - \phi'(x))(f_1(x) - cf_0(x))dx &= \int_{S^+ \cup S^-} (\phi(x) - \phi'(x))(f_1(x) - cf_0(x))dx, \text{ since } 0 \text{ when } \phi(x) = \phi'(x) \\ \int_{S^+ \cup S^-} (\phi(x) - \phi'(x))(f_1(x) - cf_0(x))dx &\geq 0, \text{ since two differences are always opposing} \\ \int_{S^+ \cup S^-} (\phi(x) - \phi'(x))f_1(x)dx &\geq \int_{S^+ \cup S^-} (\phi(x) - \phi'(x))cf_0(x)dx \geq 0, \text{ since } RHS = c[\alpha - \alpha'] \geq 0 \\ \int_{S^+ \cup S^-} \phi(x)f_1(x)dx &\geq \int_{S^+ \cup S^-} \phi'(x)f_1(x)dx \iff P_{H_1}(R) \geq P_{H_1}(R') \end{aligned}$$

7.3 Uniformly Most powerful test (UMP)

The **Most Powerful** test is the test which maximizes power under simple hypotheses, H_0, H_1 . The Neyman-Pearson Lemma tells us that the MP level- α test is the likelihood ratio test.

The **Universally Most Powerfull** test is the test that which maximizes power under composite hypotheses, H_1 . That is, for $H_1 : \theta > a$ composite, the test is MP level- α for all simple $\tilde{H}_1 \in H_1$.

The general process for showing UMP is

- Consider simple hypotheses, H_0 vs. \tilde{H}_1
- Apply the Neyman-Pearson Lemma to find MP test for H_0 vs. \tilde{H}_1
- Show that the test doesn't depend on the choice $\theta_i \in H_1$

7.3.1 Example LR and UMP test

$X_1, \dots, X_n \stackrel{i.i.d}{\sim} \text{Poisson}(\lambda), H_0 : \lambda = 1, H_1 : \lambda > 1$

$$LR(X) = \frac{\prod_{i=1}^n \exp(-\lambda_1) * \frac{\lambda_1^{X_i}}{X_i!}}{\prod_{i=1}^n \exp(-1) * \frac{1^{X_i}}{X_i!}} = \frac{\exp(-n\lambda_1) \lambda_1^{\sum X_i}}{\exp(-n)} = \exp(n(1 - \lambda_1)) \lambda_1^{\sum X_i}, \text{ choosing some } \lambda_1 \in H_1$$

$$LR(X) \geq c \iff \exp(n(1 - \lambda_1)) \lambda_1^{\sum X_i} \geq c \iff \lambda_1^{\sum X_i} \geq c' \iff \sum_{i=1}^n X_i \geq c''' = c$$

Under H_0 , $\sum_{i=1}^n X_i \sim \text{Poisson}(n)$ and level- α test rejects when $\sum_{i=1}^n X_i \geq C_{n,1-\alpha}$ (upper $(1 - \alpha)$ quantile of Poisson(n))

7.4 P-values

P-values answer "what is the smallest α that we would still reject H_0 ". For $T(X)$, a test statistic, and t , the statistic calculated from the data. Assume $T(X) \sim f_0(x)$, then $P_{H_0}[T(X) \geq t] = 1 - F_0(t) \iff \text{pval} = 1 - F_0(T(X))$

In the case $T(X) = \sqrt{n} \frac{\bar{X}_n}{\sigma} \sim N(0, \sigma)$ under H_0 , then we have $P_{H_0}[\sqrt{n} \frac{\bar{X}_n}{\sigma} \geq t] = 1 - \Phi(t) \iff \text{pval} = 1 - \Phi\left(\sqrt{n} \frac{\bar{X}_n}{\sigma}\right)$

The **distribution** of a pvalue can be described with

$$\text{pval} = 1 - F_0(T(X)) \iff P(1 - F_0(T(X)) \leq t) \iff P(T(X) \geq F_0^{-1}(1 - t)), \text{ watching signs of inequality}$$

7.5 Generalized Likelihood Ratio test

The **Generalized Likelihood Ratio test** provides us a way to compare composite hypotheses.

$$R_n = \frac{\max_{\theta \in \Omega_0 \cup \Omega_1} L_n(\theta)}{\max_{\theta \in \Omega_0} L_n(\theta)} = \frac{\hat{\theta}_{MLE}}{\max_{\theta \in \Omega_0} L_n(\theta)}$$

Twice the log of the Generalized Likelihood Ratio follows a χ_d^2 distribution with $d = k - k_0$ degrees of freedom

$$2 \log(R_n) \sim \chi_d^2, \text{ with } d = k - k_0$$

7.5.1 GLR example I

$$X_1, \dots, X_n \stackrel{i.i.d}{\sim} N(\theta, \sigma^2), H_0 : \theta = 0, H_1 : \theta \neq 0$$

$$R_n = \frac{L_n(\bar{X}_n)}{L_n(0)} = \exp\left(\frac{n\bar{X}_n^2}{2\sigma^2}\right)$$

$$2 \log(R_n) = \frac{n\bar{X}_n^2}{\sigma^2} = Z^2 \sim \chi_1^2, \text{ where } Z \sim N(0, 1)$$

7.5.2 GLR example II

The **Poisson Dispersion Test**

$$X_1, \dots, X_n \stackrel{i.i.d}{\sim} \text{Poisson}(\lambda_i), H_0 : \lambda_1 = \dots = \lambda_n, H_1 : \text{not } \lambda_i \text{ all equal}$$

$$R_n = \frac{L_n(\hat{\lambda}_{MLE_1}, \dots, \hat{\lambda}_{MLE_n})}{L_n(\bar{X}_n)} = \prod_{i=1}^n \left(\frac{X_i}{\bar{X}_n}\right)^{X_i}$$

$$2 \log(R_n) = 2 \sum_{i=1}^n X_i \log\left(\frac{X_i}{\bar{X}_n}\right) \sim \chi_{n-1}^2$$

$$2 \log(R_n) \approx \sum_{i=1}^n \frac{(X_i - \bar{X}_n)^2}{\bar{X}_n}, \text{ using Taylor approximations}$$

7.5.3 Testing multinomial distributions

We can constructing the generalized likelihood ratio in the multinomial models as well

$$X_1, \dots, X_n \sim \text{multi}(n, p_1, \dots, p_n), H_0 : p_j = p_j(\theta)$$

$$\text{Unrestrained MLE: } \hat{p}_j = \frac{X_j}{n}, \text{ MLE under } H_0: \hat{\theta}_{MLE}$$

$$R_n = \frac{\frac{n!}{X_1! \dots X_n!} \hat{p}_1^{X_1} \dots \hat{p}_n^{X_n}}{\frac{n!}{X_1! \dots X_n!} p_1(\hat{\theta})^{X_1} \dots p_n(\hat{\theta})^{X_n}} = \prod_{j=1}^n \left(\frac{\hat{p}_j}{p_j(\hat{\theta})}\right)^{X_j}$$

$$2 \log(R_n) = 2 \sum_{j=1}^n X_j \log\left(\frac{\hat{p}_j}{p_j(\hat{\theta})}\right) = 2 \sum_{j=1}^n X_j \log\left(\frac{X_j}{np_j(\hat{\theta})}\right)$$

$$2 \log(R_n) = 2 \sum_{j=1}^n O_j \log\left(\frac{O_j}{E_j}\right), \text{ for } O_j = X_j \text{ and } E_j = np_j(\hat{\theta})$$

We can approximate this equality in Generalized Likelihood Ratios using the Taylor approximation to get the **Chi Squared Statistic**

$$2 \log(R_n) = 2 \sum_{j=1}^n O_j \log\left(\frac{O_j}{E_j}\right) \approx \sum_{j=1}^n \frac{(O_j - E_j)^2}{E_j} \sim \chi_d^2, \text{ with } d = k - k_0$$

In general, we can determine the degrees of freedom under H_0 , k_0 , as $(r-1) + (c-1)$ and under H_1 , k , as $r * c - 1$. The **Chi-square Test of Homogeneity** tests $H_0 : \pi_{i1} = \dots = \pi_{iJ}$ with statistic

$$X^2 = \sum_{i=1}^I \sum_{j=1}^J \frac{(O_{ij} - E_{ij})^2}{E_{ij}} \sim \chi_{(I-1)(J-1)}^2$$

8 Helpful applied methods

8.1 Derivation of Linear regression

TODO

8.2 Confidence intervals

$$\begin{aligned} \sqrt{n} \frac{\bar{X}_n - \mu}{\sigma} &\xrightarrow{d} Z \sim N(0, 1), \text{ by CLT for calculated } \bar{X}_n \\ \sqrt{n} \frac{\bar{X}_n - \mu}{\hat{\sigma}} &\xrightarrow{d} Z \sim N(0, 1), \text{ by Slutsky's theorem} \\ P(Z \leq \Phi(1 - \alpha)) &= P(Z \leq Z_{1-\alpha}) = 1 - \alpha \iff P(Z_{\alpha/2} \leq Z \leq Z_{1-\alpha/2}) = 1 - \alpha \\ P(Z_{\alpha/2} \leq \sqrt{n} \frac{\bar{X}_n - \mu}{\hat{\sigma}} \leq Z_{1-\alpha/2}) &= P\left(\frac{\hat{\sigma}}{\sqrt{n}} Z_{\alpha/2} \leq \bar{X}_n - \mu \leq \frac{\hat{\sigma}}{\sqrt{n}} Z_{1-\alpha/2}\right) \\ &= P\left(\bar{X}_n - \frac{\hat{\sigma}}{\sqrt{n}} Z_{1-\alpha/2} \leq \mu \leq \bar{X}_n + \frac{\hat{\sigma}}{\sqrt{n}} Z_{1-\alpha/2}\right) \xrightarrow{d} 1 - \alpha \\ \mu &\in \left[\bar{X}_n \pm \frac{\hat{\sigma}}{\sqrt{n}} Z_{1-\alpha/2} \right] \text{ with } p \xrightarrow{d} 1 - \alpha \end{aligned}$$

8.3 Asymptotic distribution of sample variance

$$\begin{aligned} S_n^2 &= \frac{1}{n-1} \sum_{i=1}^n (X_i - \bar{X}_n)^2 = \frac{n}{n-1} \left[\frac{1}{n} \sum_{i=1}^n (X_i - \mu)^2 - (\bar{X}_n - \mu)^2 \right] \\ \sqrt{n}(S_n^2 - \sigma^2) &= \frac{n}{n-1} \left[\sqrt{n} \left(\frac{1}{n} \sum_{i=1}^n (X_i - \mu)^2 \right) + \sqrt{n}(\bar{X}_n - \mu)^2 \right] - \sqrt{n}\sigma^2 \\ &= \frac{n}{n-1} \left[\sqrt{n} \left(\frac{1}{n} \sum_{i=1}^n (X_i - \mu)^2 - \sigma^2 \right) + \sqrt{n}\sigma^2 + \sqrt{n}(\bar{X}_n - \mu)^2 \right] \\ &= \frac{n}{n-1} * \sqrt{n} \left(\frac{1}{n} \sum_{i=1}^n (X_i - \mu)^2 - \sigma^2 \right) + \frac{\sqrt{n}}{n-1} \sigma^2 + \frac{n\sqrt{n}}{n-1} \sqrt{n}(\bar{X}_n - \mu)^2 \\ \frac{n}{n-1} &\xrightarrow{p} 1 \bullet \frac{\sqrt{n}\sigma^2}{n-1} \xrightarrow{p} 0 \bullet \sqrt{n}(\bar{X}_n - \mu)^2 \xrightarrow{p} 0, \text{ since by Slutsky } (\bar{X}_n - \mu) \xrightarrow{p} 0 \text{ \& } \sqrt{n}(\bar{X}_n - \mu) \xrightarrow{d} N(0, 1) \\ \therefore \sqrt{n}(S_n^2 - \sigma^2) &\xrightarrow{d} \sqrt{n} \left(\frac{1}{n} \sum_{i=1}^n (X_i - \mu)^2 - \sigma^2 \right) \xrightarrow{d} N(0, \text{Var}[(X_i - \mu)^2]), \text{ since } E \left[\frac{1}{n} \sum_{i=1}^n (X_i - \mu)^2 \right] = \sigma^2 \end{aligned}$$

8.4 Comparing two samples

TODO

8.5 Fischer's Exact test

TODO

TODO – You're on HW 7